

- [17] John F Dovidio and Samuel L Gaertner. Intergroup bias. 2010.
- [18] Lilliana Mason and Julie Wronski. One tribe to bind them all: How our social group attachments strengthen partisanship. *Political Psychology*, 39:257–277, 2018.
- [19] Lilliana Mason. A cross-cutting calm: How social sorting drives affective polarization. *Public Opinion Quarterly*, 80(S1):351–377, 2016.
- [20] Yukun Zhao, Zhen Huang, Martin Seligman, and Kaiping Peng. Risk and prosocial behavioural cues elicit human-like response patterns from ai chatbots. *Scientific reports*, 14(1):7095, 2024.
- [21] Thilo Hagendorff. Deception abilities emerged in large language models. *Proceedings of the National Academy of Sciences*, 121(24):e2317967121, 2024.
- [22] Lindia Tjuatja, Valerie Chen, Sherry Tongshuang Wu, Ameet Talwalkar, and Graham Neubig. Do llms exhibit human-like response biases? a case study in survey design. *arXiv preprint arXiv:2311.04076*, 2023.
- [23] Tiancheng Hu, Yara Kyrychenko, Steve Rathje, Nigel Collier, Sander van der Linden, and Jon Roozenbeek. Generative language models exhibit social identity biases. *arXiv preprint arXiv:2310.15819*, 2023.
- [24] Ruibo Liu, Chenyan Jia, Jason Wei, Guangxuan Xu, Lili Wang, and Soroush Vosoughi. Mitigating political bias in language models through reinforced calibration. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 14857–14866, 2021.
- [25] David Rozado. The political biases of chatgpt. *Social Sciences*, 12(3):148, 2023.
- [26] SunYoung Park, Kyuri Choi, Haeun Yu, and Youngjoong Ko. Never too late to learn: Regularizing gender bias in coreference resolution. In *Web Search and Data Mining, WSDM ’23*, pages 15–23, 2023.
- [27] Xiao Fang, Shangkun Che, Minjia Mao, Hongzhe Zhang, Ming Zhao, and Xiaohang Zhao. Bias of ai-generated content: an examination of news produced by large language models. *Scientific Reports*, 14(1):5224, 2024.
- [28] Rishav Hada, Agrima Seth, Harshita Diddee, and Kalika Bali. “fifty shades of bias”: Normative ratings of gender bias in GPT generated English text. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 1862–1876. Association for Computational Linguistics, December 2023.
- [29] Myra Cheng, Tiziano Piccardi, and Diyi Yang. Compost: Characterizing and evaluating caricature in llm simulations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 10853–10875, 2023.
- [30] Andy Liu, Mona Diab, and Daniel Fried. Evaluating large language model biases in persona-steered generation. *arXiv preprint arXiv:2405.20253*, 2024.
- [31] Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, et al. A survey on evaluation of large language models. *ACM Transactions on Intelligent Systems and Technology*, 15(3):1–45, 2024.
- [32] Fabrizio Gilardi, Meysam Alizadeh, and Maël Kubli. Chatgpt outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences*, 120(30):e2305016120, 2023.
- [33] Gati V Aher, Rosa I Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. In *International Conference on Machine Learning*, pages 337–371. PMLR, 2023.
- [34] Isabelle Augenstein, Timothy Baldwin, Meeyoung Cha, Tanmoy Chakraborty, Giovanni Luca Ciampaglia, David Corney, Renee DiResta, Emilio Ferrara, Scott Hale, Alon Halevy, et al. Factuality challenges in the era of large language models. *arXiv preprint arXiv:2310.05189*, 2023.
- [35] Xiaosu Wang, Yun Xiong, Beichen Kang, Yao Zhang, Philip S Yu, and Yangyong Zhu. Reducing negative effects of the biases of language models in zero-shot setting. In *Web Search and Data Mining, WSDM ’23*, pages 904–912, 2023.
- [36] Jessica Echterhoff, Yao Liu, Abeer Alessa, Julian McAuley, and Zexue He. Cognitive bias in high-stakes decision-making with llms. *arXiv preprint arXiv:2403.00811*, 2024.
- [37] Shaz Furniturewala, Surgan Jandial, Abhinav Java, Pragyan Banerjee, Simra Shahid, Sumit Bhatia, and Kokil Jaidka. Thinking fair and slow: On the efficacy of structured prompts for debiasing language models. *arXiv preprint arXiv:2405.10431*, 2024.

- [38] Dominic Ed Abrams and Michael A Hogg. *Social identity theory: Constructive and critical advances*. Springer-Verlag Publishing, 1990.
- [39] Michael A Hogg. *Social identity theory*. Springer, 2016.
- [40] Nicole Tausch, Katharina Schmid, and Miles Hewstone. The social psychology of intergroup relations. *Handbook on peace education*, pages 75–86, 2010.
- [41] Henri Tajfel. Experiments in intergroup discrimination. *Scientific american*, 223(5):96–103, 1970.
- [42] Henri Tajfel, Michael G Billig, Robert P Bundy, and Claude Flament. Social categorization and intergroup behaviour. *European journal of social psychology*, 1(2):149–178, 1971.
- [43] John C Turner. Social comparison and social identity: Some prospects for intergroup behaviour. *European journal of social psychology*, 5(1):1–34, 1975.
- [44] Amelie Mummendey and Michael Wenzel. Social discrimination and tolerance in intergroup relations: Reactions to intergroup difference. *Personality and social psychology review*, 3(2):158–174, 1999.
- [45] Fabio Motoki, Valdemar Pinho Neto, and Victor Rodrigues. More human than human: measuring chatgpt political bias. *Public Choice*, 198(1):3–23, 2024.
- [46] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35:22199–22213, 2022.
- [47] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [48] Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023.
- [49] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- [50] Dominic Abrams, Margaret Wetherell, Sandra Cochrane, Michael A Hogg, and John C Turner. Knowing what to think by knowing who you are: Self-categorization and the nature of norm formation, conformity and group polarization. *British journal of social psychology*, 29(2):97–119, 1990.
- [51] Michael A Hogg, John C Turner, and Barbara Davidson. Polarized norms and social frames of reference: A test of the self-categorization theory of group polarization. *Basic and Applied Social Psychology*, 11(1):77–100, 1990.
- [52] Michael A Hogg. Social identity. *Handbook of self and identity*, pages 462–479, 2003.
- [53] Diane Mackie and Joel Cooper. Attitude polarization: Effects of group membership. *Journal of Personality and Social Psychology*, 46(3):575, 1984.
- [54] Diane M Mackie. Social identification effects in group polarization. *Journal of Personality and Social psychology*, 50(4):720, 1986.
- [55] Jeff Greenberg and Eva Jonas. Psychological motives and political orientation—the left, the right, and the rigid: comment on jost et al. 2003.
- [56] Merriam-Webster. What is the difference between a democrat and a republican?, 2023.
- [57] Scott Clifford. Compassionate democrats and tough republicans: How ideology shapes partisan stereotypes. *Political Behavior*, 42(4):1269–1293, 2020.
- [58] Evan Weingarten, Qijia Chen, Maxwell McAdams, Jessica Yi, Justin Hepler, and Dolores Albarracín. From primed concepts to action: A meta-analysis of the behavioral effects of incidentally presented words. *Psychological bulletin*, 142(5):472, 2016.
- [59] Peter Glick and Susan T Fiske. The ambivalence toward men inventory: Differentiating hostile and benevolent beliefs about men. *Psychology of women quarterly*, 23(3):519–536, 1999.
- [60] Peter Glick and Susan T Fiske. The ambivalent sexism inventory: Differentiating hostile and benevolent sexism. In *Social cognition*, pages 116–160. Routledge, 2018.

- [61] Gordon Willard Allport, Kenneth Clark, and Thomas Pettigrew. The nature of prejudice. 1954.
- [62] Wei Wei, Xubin Ren, Jiabin Tang, Qinyong Wang, Lixin Su, Suqi Cheng, Junfeng Wang, Dawei Yin, and Chao Huang. Llmrec: Large language models with graph augmentation for recommendation. In *Web Search and Data Mining*, WSDM '24, page 806–815, 2024.
- [63] Qijiong Liu, Nuo Chen, Tetsuya Sakai, and Xiao-Ming Wu. Once: Boosting content-based recommendation with both open-and closed-source large language models. In *Web Search and Data Mining*, WSDM '24, pages 452–461, 2024.
- [64] Maurice Jakesch, Jeffrey T Hancock, and Mor Naaman. Human heuristics for ai-generated language are flawed. *Proceedings of the National Academy of Sciences*, 120(11):e2208839120, 2023.
- [65] Paul Rozin and Edward B Royzman. Negativity bias, negativity dominance, and contagion. *Personality and social psychology review*, 5(4):296–320, 2001.
- [66] Emilio Ferrara. Social bot detection in the age of chatgpt: Challenges and opportunities. *First Monday*, 2023.
- [67] Kobi Hackenburg and Helen Margetts. Evaluating the persuasive influence of political microtargeting with large language models. *Proceedings of the National Academy of Sciences*, 121(24):e2403116121, 2024.
- [68] Rishav Hada, Amir Ebrahimi Fard, Sarah Shugars, Federico Bianchi, Patricia Rossini, Dirk Hovy, Rebekah Tromble, and Nava Tintarev. Beyond digital" echo chambers": The role of viewpoint diversity in political discussion. In *Web Search and Data Mining*, WSDM '23, pages 33–41, 2023.
- [69] Uthsav Chitra and Christopher Musco. Analyzing the impact of filter bubbles on social network polarization. In *Web Search and Data Mining*, WSDM '20, pages 115–123, 2020.
- [70] Chengcheng Shao, Giovanni Luca Ciampaglia, Onur Varol, Kai-Cheng Yang, Alessandro Flammini, and Filippo Menczer. The spread of low-credibility content by social bots. *Nature communications*, 9(1):1–9, 2018.
- [71] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. Ai models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759, 2024.
- [72] David G Rand, Thomas Pfeiffer, Anna Dreber, Rachel W Sheketoff, Nils C Wernerfelt, and Yochai Benkler. Dynamic remodeling of in-group bias during the 2008 presidential election. *Proceedings of the National Academy of Sciences*, 106(15):6187–6191, 2009.
- [73] Jonas Wallat, Adam Jatowt, and Avishek Anand. Temporal blind spots in large language models. In *Web Search and Data Mining*, WSDM '24, page 683–692, 2024.
- [74] Yuyan Chen, Jin Zhao, Zhihao Wen, Zhixu Li, and Yanghua Xiao. Temporalmed: Advancing medical dialogues with time-aware responses in large language models. In *Web Search and Data Mining*, WSDM '24, pages 116–124, 2024.
- [75] Bingbing Wen, Jihan Yao, Shangbin Feng, Chenjun Xu, Yulia Tsvetkov, Bill Howe, and Lucy Lu Wang. The art of refusal: A survey of abstention in large language models. *arXiv preprint arXiv:2407.18418*, 2024.